

Lecture 4 Notebook

September 4, 2018

1 Lecture 4 Notebook

Duncan Callaway September 4 2018

This lecture continues introducing the class to Pandas and goes into "groupby"

In class I worked in "Duncan's Lecture 4 in class workbook.ipynb"

```
In [ ]: import numpy as np
import pandas as pd
```

1.1 Recap last lecture

Data frame vs dict of lists

```
In [ ]: fruit_info={'fruit':['apple','banana','orange','raspberry'],
                    'color':['red','yellow','orange','pink'],
                    'weight':[120,150,250,15]}
fruit_info_df = pd.DataFrame(data = fruit_info)
print(fruit_info)
fruit_info_df
```

The data frame has 1. column headers, 2. An index column 2. rows 4. columns 5. numeric and text entries -- but columns are all the same type.

Note, last time I tried getting the "index" value from the fruit info dataframe:

"the index" Note this is different from locational indexing. We're talking about the column that identifies the row of the frame.

```
In [ ]: fruit_info_df.index
```

What I was *expecting* was a list from 0 to 3, e.g. [0, 1, 2, 3]

But I got the above. This is just an alternative way of giving the same information. However note, if I do this:

```
In [ ]: fruit_info_df.index = ['zero', 'one', 'two', 'three']
```

```
In [ ]: fruit_info_df.index
```

...then I get what I expected. More on indices in a moment.

loc and iloc loc identifies location by column and header names.

```
In [ ]: fruit_info_df.loc['zero':'two', 'fruit':'weight']
```

note, loc is inclusive!

.iloc identifies location by number -- just like indexing in numpy.

```
In [ ]: fruit_info_df.iloc[0:3,1:3]
```

.iloc is exclusive on the end location value.

1.2 Back to our question: which hour had the most wind...

```
In [ ]: caiso_data_stack = pd.read_csv('CAISO_2017to2018_stack.csv', index_col= 0)
```

Let's make the name shorter to save me typing:

```
In [ ]: cds = caiso_data_stack
        cds.head()
```

Let's look at some info about the data:

```
In [ ]: cds.shape
```

```
In [ ]: cds.size
```

What did those two commands give us? shape: (number of rows, number of columns of *data*)
size: total number of cell entries.

Notice these numbers don't include what's in the index.

Here's something fun --

```
In [ ]: cds.describe()
```

1.3 Logical indexing

Logical indexing is an extremely powerful way to pull data out of a frame.

For example, with the stacked data frame, let's pull out only wind generation.

First, I'll show you a boolean series based on comparisons to the 'Source' data column:

```
In [ ]: (cds['Source']=='WIND TOTAL').head()
```

Now we can embed that inside the .loc method:

```
In [ ]: cds.loc[cds['Source']=='WIND TOTAL',:].head()
```

Ok. Any ideas how we can use that to get the information we want? Reminder, the question is:

What hour of the day had the lowest average wind power in California in the last 12 months?

```
In [ ]: wind = cds.loc[cds['Source']=='WIND TOTAL',:]
```

What is the data structure of wind?

```
In [ ]: type(wind)
```

Next week we'll use pivots to do this better, but for now let's use a for loop to get information by hour.

First thing to do is figure out how to get the hour out of the index.

`datetime.strptime` is useful for this if you're working on individual dates.

But `pd.to_datetime` is even better, especially if you're working on a lot of values in a list (or as the case will be, values in a pandas series).

```
In [ ]: windindex = pd.to_datetime(wind.index)
        windindex.hour
```

```
In [ ]: wind_ave = [] # initializes a list to populate
        for i in range(0,24):
            wind_ave.append(np.mean(wind.loc[windindex.hour == i,:]))
```

```
In [ ]: print(wind_ave)
```

```
In [ ]: type(wind_ave)
```

```
In [ ]: import matplotlib.pyplot as plt
```

```
In [ ]: plt.plot(wind_ave)
```

We can see pretty clearly that the min is 10 or 11...let's dig a little more.

One way to do this is to drop the data into a data frame and then *sort* the data frame.

```
In [ ]: df_wind = pd.DataFrame(wind_ave)
        df_wind
```

I'm going to be adding more MWh values to the data frame in just a moment, so let's be clear that this is the average

```
In [ ]: df_wind.columns = ['Average MWh']
```

```
In [ ]: df_wind.sort_values(by='Average MWh',ascending=True).head()
```

Ok -- so it looks as though mid-day is the minimum *average*.

Nice to see that the index values were preserved

But what's the range?

```
In [ ]: wind_min = [] # initializes a list to populate
        wind_max = [] # initializes a list to populate
        for i in range(0,24):
            wind_min.append(np.min(wind.loc[windindex.hour == i,:]))
            wind_max.append(np.max(wind.loc[windindex.hour == i,:]))
```

```
In [ ]: wind_max[0]
```

```
In [ ]: df_wind['min MWh']=pd.DataFrame(wind_min)['MWh']
        df_wind['max MWh']=pd.DataFrame(wind_max)['MWh']
```

```
In [ ]: df_wind
```

```
In [ ]: plt.plot(df_wind)
```

1.4 Row and column labels

The columns are identified with a list of values. Let's look at the fruit data set again:

```
In [ ]: fruit_info_df.columns
```

```
In [ ]: type(fruit_info_df.columns)
```

The rows are similarly labeled:

```
In [ ]: fruit_info_df.index
```

```
In [ ]: type(fruit_info_df.index)
```

They are both the same data type within Pandas -- the "Index"

Note, we can do a bunch of other stuff:

1.5 Merging

Lets make another data frame and tack it on to the first

```
In [ ]: price_df = pd.DataFrame({'price':[0.5, 0.65, 1, 0.15],  
                                'frut':['apple', 'banana', 'orange', 'rasberry']})  
price_df
```

```
In [ ]: fruit_info_df
```

```
In [ ]: pd.merge(price_df,fruit_info_df)
```

What went wrong?

First, we didn't spell fruit correctly. Two ways to fix. First, specify the columns directly:

```
In [ ]: pd.merge(price_df,fruit_info_df, left_on = 'frut', right_on = 'fruit')
```

Second, fix the spelling and *don't* tell pandas. In this case pandas works to figure out what's in common.

```
In [ ]: price_df.columns[0]='fruit'
```

Bummer! Can't mutate index values. What to do?

```
In [ ]: col_list = list(price_df.columns)  
col_list
```

```
In [ ]: col_list[0] = 'fruit'
```

```
In [ ]: price_df.columns = col_list  
price_df
```

```
In [ ]: pd.merge(fruit_info_df,price_df)
```

Note we can use different syntax:

```
In [ ]: fruit_info_df.merge(price_df)
```

Now we're still missing raspberries -- why?
Again, spelling error in the new frame. Let's fix:

```
In [ ]: price_df.loc[3,'fruit'] = 'raspberry'
```

Note we could change individual entries in the data frame itself. They are mutable.

```
In [ ]: fruit_info_df.merge(price_df)
```

Note the fruit_info data frame is still intact, you'd need to assign it to a data frame name to save it.

```
In [ ]: fruit_info_df
```

Here's a cool little factoid about data frames: you can write for loops that burn through the columns of the frame.

```
In [ ]: for i in fruit_info_df:
        print(fruit_info_df.loc['one',i])
```

Note, there are other commands -- join, concat, and these do similar things.

I haven't learned enough to carefully choose between them, but merge seems to work well.

FWIW, pd.concat seems to be a little more brute force -- requires more careful syntax, but likely does unexpected things less often once you understand the syntax.

```
In [ ]: pd.concat([fruit_info_df,price_df])
```

You can see in the above that setting axis equal to 1 will join by appending columns rather than rows.

```
In [ ]: merged_df = fruit_info_df.merge(price_df)
        merged_df
```

We can streamline by replacing the index number with the fruit column.

What's the inplace command for? It means the re-defined dataframe is assigned to the original name. This is advantageous in memory constrained situations.

```
In [ ]: merged_df.set_index('fruit', inplace = True)
        merged_df
```

1.6 Multilevel indexing

We can also assign "multilevel" column or row names, like so:

```
In [ ]: levels = [('categorical', 'color'),('quantitative', 'weight'),('quantitative', 'price')]
        levels
```

Note the use of tuples (sets of values in parentheses) in setting up multiindex. This will come again later.

```
In [ ]: merged_df.columns = pd.MultiIndex.from_tuples(levels)
        merged_df
```

Now we have categories and subcategories of columns:

```
In [ ]: merged_df['quantitative']
```

Note, we can also drop and add things. With multilevel indexing things get a little tricky. First, we can drop everything from the top level:

```
In [ ]: merged_test_df = merged_df.drop(columns=[('quantitative',)], axis = 1)
        merged_test_df
```

Note that I put the column identifier inside the parens, like a tuple, but it's not essential there.

However if we want to drop only a column from the second level, we get an error without the tuple syntax:

```
In [ ]: merged_test_df = merged_df.drop(columns=[('quantitative', 'price')], axis = 1)
        merged_test_df
```

We can also drop rows:

```
In [ ]: merged_df.drop(index=[('apple')], axis = 0, inplace = True)
        merged_df
```

Note indexing multilevels with `.loc` gets a little tricky. The thing to keep in mind is that you're working with tuples in each index location:

```
In [ ]: merged_df.loc['banana', ('quantitative', 'price')]
```

If you leave an entry of the tuple empty you get all values.

```
In [ ]: merged_df.loc['banana', ('quantitative', )]
```

You can also loop through the columns of the multilevel data frame like this:

```
In [ ]: for i, j in merged_df:
        print(merged_df.loc['banana', (i, j)])
```

1.7 Groupby

(these notes adapted from last Spring's DS100 notebook)

Let's make a toy DF (example taken from Wes McKinney's [Python for Data Analysis](#)):

```
In [ ]: df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                           'key2' : ['one', 'two', 'one', 'two', 'one'],
                           'data1' : np.random.randn(5),
                           'data2' : np.random.randn(5)})

df
```

Let's group just the data1 column by the key1 column. A call to `groupby` does that.

Note, the syntax is to begin by invoking the portion of the dataframe we want to group (here, `df['data1']`), then we apply the `groupby` method with the portion of the dataframe we want to group on (here `df['key1']`)

What is the object that results?

```
In [ ]: grouped = df['data1'].groupby(df['key1'])
        grouped
```

As we see, it's not simply a new DataFrame. Instead, it's an object, in this case `SeriesGroupBy`. We'll see in a moment that if we group many columns of data we get a `DataFrameGroupBy` object.

To look inside we need to use different syntax. The specific thing we're looking for are the groups of the object...but let's tab in to the `grouped` object to see what's there.

```
In [ ]: grouped.groups
```

That gave us the groups (a and b) and the indices of elements in the groups, but nothing else.

If we call `grouped.groups` elements, we don't get much of use; we wind up just retrieving the elements of the list above:

```
In [ ]: grouped.groups['a'][2]
```

But the `grouped` object is capable of making computations across all groups -- this is where it gets powerful.

We can try things like `.count()`, `.min()` and `.mean()`.

Notice if you don't put the parens after the method, pandas returns information about what the method does, but not its actual output.

```
In [ ]: grouped.count()
```

But it can be informative to look at what's inside. We can iterate over a `groupby` object, as we iterate we get pairs of (name, group), where the group is either a `Series` or a `DataFrame`, depending on whether the `groupby` object is a `SeriesGroupBy` (as above) or a `DataFrameGroupBy` (see below):

```
In [ ]: from IPython.display import display # like print, but for complex objects

        for name, group in grouped:
            print('Name:', name)
            display(group)
```

We can group on multiple keys, and the result is grouping by tuples:

```
In [ ]: g2 = df['data1'].groupby([df['key1'], df['key2']])
        g2
```

```
In [ ]: g2.groups
```

Let's look at the dataframe again, for a reminder:

```
In [ ]: df
```

```
In [ ]: g2.mean()
```

We can also group the entire dataframe -- not just one column of it -- on a single key. This results in a `DataFrameGroupBy` object as the result:

```
In [ ]: k1g = df.groupby('key1')
        k1g
```

```
In [ ]: k1g.groups
```

```
In [ ]: k1g.mean()
```

But let's look at what's inside of `k1g`:

```
In [ ]: for n, g in k1g:
        print('name:', n)
        display(g)
```

Where did column `key2` go in the mean above? It's a *nuisance column*, which gets automatically eliminated from an operation where it doesn't make sense (such as a numerical mean).

1.7.1 Grouping over a different dimension

Above, we've been grouping data along the rows, using column keys as our selectors.

But we can also group along the *columns*,
What's even more cool? We can group by *data type*.
Here we'll group along columns, by data type:

```
In [ ]: df.dtypes
```

```
In [ ]: grouped = df.groupby(df.dtypes, axis=1)
        for dtype, group in grouped:
            print(dtype)
            display(group)
```

1.8 Let's take the quiz.

1.9 Using `groupby` to re-ask our question

Which hour had the lowest average wind production?

```
In [ ]: cds.head()
```

It will help to have a column of hour of day values:

```
In [ ]: cds_time = pd.to_datetime(cds.index)
        cds_time.hour
```

Let's add that list of values into the data frame.


```
In [ ]: cds['hour'] = cds_time.hour
```

```
In [ ]: cds.head(10)
```

Now do the grouping.

See if you can do it yourself: we want to group MWh values by source AND hour.

```
In [ ]: cds_grouped = cds['MWh'].groupby([cds['Source'],cds['hour']])
```

```
In [ ]: cds_grouped.groups
```

Now we can see *all* the means for all sources and hours.

Didn't need to do any fancy logical indexing or looping!

```
In [ ]: cds_grouped.mean()
```

Now it would be nice to see that information in a dataframe, wouldn't it?

```
In [ ]: averages = pd.DataFrame(cds_grouped.mean())
```

```
In [ ]: averages
```

And lo and behold, we have a multilevel index for the rows!

```
In [ ]: averages.loc[('WIND TOTAL',),:]
```

But now we can look at other sources

```
In [ ]: averages.index
```

```
In [ ]: averages.loc[('SMALL HYDRO',),:]
```

```
In [ ]: plt.plot(averages.loc[('SMALL HYDRO',),:])
```

```
In [ ]: plt.plot(averages.loc[('GEOTHERMAL',),:])
```

```
In [ ]: plt.plot(averages.loc[('SOLAR PV',),:])
```

```
In [ ]: plt.plot(cds.loc[cds.loc[:, 'Source']=='SOLAR PV', 'MWh'])
```