# Data, Environment and Society: Lecture 22: Measuring Classifier Performance

Instructor: Duncan Callaway

GSI: Salma Elmallah

November 14, 2019

#### **Announcements**

- Today:
  - Quick discussion of classifier performance, with notebook
  - Exam review
- HW10 posts today, due Tuesday before Thansksgiving (11/26)
- Lab next week: Cancelled, prep for exam! Salma will hold special office hours.
- Tuesday next week: Exam
  - Covers up to Lecture 19 21
  - Covers up to HW9

## Objectives for today

- Methods for measuring performance of classifiers
- Reminders on parameter, hyperparameter tuning

Reminder: classification error rate

## Reminder: classification error rate

The typical error, RSS  $=\sum_{i=1}^{N}(y_i-\hat{y}_i)^2$  won't work.

Alternatives? Let's start by defining

 $p_{mk} =$  fraction of observations belonging to class k in region m.

Then a simple measure is:

Classification error rate = how many training observations don't fall into the assigned class.

Within-region this is simply:

$$E_m = 1 - \max_k(\hat{p}_{mk})$$

### The confusion matrix

- A "confusion matrix" C is such that  $C_{i,j}$  is equal to the number of observations known to be in group i but predicted to be in group j.
- Numbers on the diagonal of the matrix are correct predictions
- Rows sum to the number of actual observations in a category
- Columns sum to the number of predicted observations in a category

		Predict: False
Actual: True	True pos (TP)	False neg (FN)
Actual: False	False pos (FP)	True neg (TN)

### Precision and recall

	Predict: True	Predict: False
Actual: True	True pos (TP)	False neg (FN)
Actual: False	False pos (FP)	True neg (TN)

Precision: Correct number of positives divided by total number of positive predictions.
This is a "true positive rate": How many things classified as positive are actually positive?

$$=\frac{TP}{TP+FP}$$

Recall: Correctly number of positives divided by total number of true positives.
This is (1- false positive rate): How many positive outcomes are classified as positive?

$$=\frac{TP}{TP+FN}$$

# Reminder / clarification for cross validation: Models have two different classes of parameters

Parameters that enter as decision variables for minimizing loss function:

Shrinkage: 
$$\beta^*$$
 =  $\arg\min_{\beta} \sum_{i=1}^{N} (Y_i - X_i \beta)^2 + \lambda \cdot R(\beta)$    
  $\operatorname{CART:} \{j^*, s^*\}$  =  $\arg\min_{j \in J, s \in X_j} \sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$ 

"Hyperparameters": parameters that constrain how you solve the loss function. These generally prevent overfit.

# Reminder / clarification for cross validation: Models have two different classes of parameters

Parameters that enter as decision variables for minimizing loss function:

Shrinkage: 
$$\beta^*$$
 =  $\arg\min_{\beta} \sum_{i=1}^{N} (Y_i - X_i \beta)^2 + \lambda \cdot R(\beta)$    
 CART:  $\{j^*, s^*\}$  =  $\arg\min_{j \in J, s \in X_j} \sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$ 

- "Hyperparameters": parameters that constrain how you solve the loss function. These generally prevent overfit.
  - $\triangleright$   $\lambda$  in shrinkage methods
  - How deep to grow a classification tree?
  - $\triangleright \sum_{i=1}^{n} \epsilon_i$  in SVM

Ways to minimize the loss function

## Ways to minimize the loss function

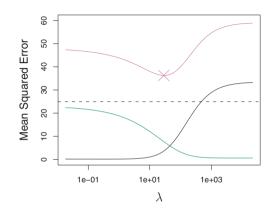
- Closed form solution e.g. normal equations
- Gradient search

In both cases, we're relying on the condition that the gradient of the loss function approaches zero as we approach the optimal solution

## Ways to choose hyperparameters

## Ways to choose hyperparameters

- **Grid search**: This is what we've done with shrinkage methods, when there is just one parameter to tune  $(\lambda)$
- Randomized parameter search: This is what you're doing in HW10. Works well when you have lots of hyperparameters to tune.



### In both cases, all we're doing is

- Creating a list of candidate hyperparameters (or sets of hyperparameters)
- Training the model (with the training data) for each hyperparameter in the list
- Choosing the hyperparameter with the best cross-validated error.