

# Data, Environment and Society:

## Lecture 25: Neural Networks

Instructor: Duncan Callaway  
GSI: Salma Elmallah

**November 26, 2019**

# Announcements

- HW10 due today
- Next today and tuesday: neural nets
  - ▶ Feel free to use on projects – but no HW here.
- Course evaluations available online; I will make time next Tuesday
- Thursday: Career panel
  - ▶ JP Dolphin, manager of Strategic Data Science at PG&E
  - ▶ Tanner Burke, senior data engineer, Streetlight Data
  - ▶ Jason Harville, assistant executive director of the Energy Data and Analytics Office, California Energy Commission

# Today's outline


## ① Neural networks (NN)

- ▶ Brief introduction
- ▶ Experiment with tensorflow playground – try fitting different classification problems
- ▶ Objective: Understand the role of key parameters, what the hyperparameters are, and the model fitting process

## ② Exam handout and discussion

# Neural networks: Origins

- The name is due to analogy with brains
- First developed in 1943
- Inspired the development of the perceptron (see HW10) in the '50s
  - ▶ Here the purpose was just to remove noise on phone lines
  - ▶ Not to reproduce thought...
- Little research activity ~1960-1990's due to computing limitations
  - ▶ Major exception: Werbos developed back-propagation in 1974. First effort to get NN to "learn" parameters
- Computing advances made "deep" NN possible in the last 20 years



# Mathematics for a single “neuron”

In words, each neuron...

- Takes a vector of values as inputs
- Creates a scalar from a linear combination of the vector entries
- Passes the resulting scalar through an “activation function”
- Outputs a single value from that activation function

Terminology analogies:

- Electrical signal from other cells  $\Leftrightarrow$  input
- Neuron  $\Leftrightarrow$  Activation function
- Electrical signal to other cells  $\Leftrightarrow$  output


# Mathematics for a single “neuron”

In words, each neuron...

- Takes a vector of values as inputs
- Creates a scalar from a linear combination of the vector entries
- Passes the resulting scalar through an “activation function”
- Outputs a single value from that activation function

Terminology analogies:

- Electrical signal from other cells  $\rightleftharpoons$  input
- Neuron  $\rightleftharpoons$  Activation function
- Electrical signal to other cells  $\rightleftharpoons$  output



## What's $f$ , the activation function?


① sigmoid

② tanh

③ rectified linear (ReLU)


# What's $f$ , the activation function?

① sigmoid




$$f(v) = \frac{1}{1 + e^{-v}}$$

② tanh



$$f(v) = \frac{e^v - e^{-v}}{e^v + e^{-v}}$$

③ rectified linear (ReLU)




$$\max(0, v)$$



## How the sigmoid function works

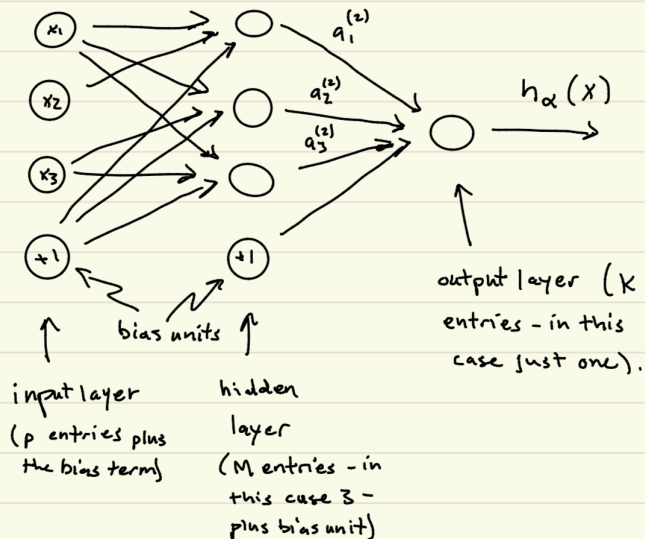
## How the sigmoid function works



Remember:  $v = \alpha_0 + \alpha^T \mathbf{x}$

Neural network: just gang the neurons together

## Neural network: just gang the neurons together



Convention:  $a_i^{(l)}$  is the output of the  $i^{\text{th}}$  neuron in the  $l^{\text{th}}$  layer.

# Mathematical merging of neurons

Convention:

- $\alpha_{ij}^{(l)} \rightarrow$  weight from node  $j$  in layer  $l$  to node  $i$  in  $l + 1$  layer.
- $a_i^{(l)} \rightarrow$  output of node  $i$  in layer  $l$ .

$$a_1^{(2)} =$$

$$a_1^{(3)} =$$

Note that I used  $x$  in the first equation because I'm calling the features (inputs to the model) the first "layer" of the network

**Question:** What are the parameters of a neural network model?

# Mathematical merging of neurons

Convention:

- $\alpha_{ij}^{(l)} \rightarrow$  weight from node  $j$  in layer  $l$  to node  $i$  in  $l + 1$  layer.
- $a_i^{(l)} \rightarrow$  output of node  $i$  in layer  $l$ .

$$a_1^{(2)} = f(\alpha_{10}^{(1)} + \alpha_{11}^{(1)} x_1^{(1)} + \alpha_{12}^{(1)} x_2^{(1)} + \alpha_{13}^{(1)} x_3^{(1)})$$

$$a_1^{(3)} = f(\alpha_{10}^{(2)} + \sum_{j=1}^M \alpha_{1j}^{(2)} a_j^{(2)})$$

Note that I used  $x$  in the first equation because I'm calling the features (inputs to the model) the first "layer" of the network

**Question:** What are the parameters of a neural network model?

Just the  $\alpha$  values.  $a$  values are outputs from internal nodes or neurons. We call these "hidden states" because they depend on the input values  $x$ .

## Thinking about the features and target

Let's watch this video. It uses graphics in a nice way to explain what NNs are doing.

<https://www.youtube.com/watch?v=aircAruvnKk>

Start the video at 2:05. We'll stop watching around 5:30.

Compact notation motivates a name...



## Compact notation motivates a name...

$$a_i^{(2)} = f(\alpha_{i0}^{(1)} + \alpha_{i1}^{(1)} x_1^{(1)} + \alpha_{i2}^{(1)} x_2^{(1)} + \alpha_{i3}^{(1)} x_3^{(1)})$$

$$a_i^{(3)} = f(\alpha_{i0}^{(2)} + \sum_{j=1}^{M_2} \alpha_{ij}^{(2)} a_j^{(2)}) \quad (M_j \text{ is the number of neurons in layer } j)$$

$$a_i^{(4)} = f(\alpha_{i0}^{(3)} + \sum_{j=1}^{M_3} \alpha_{ij}^{(3)} a_j^{(3)})$$

$\vdots$

$$h_\alpha(x) = f(a^{(\ell)}, \alpha^{(\ell)}) \quad \text{Final output of NN. } \ell \text{ is the number of layers}$$

- The  $\alpha$  subscript means  $h$  is a function of ALL the  $\alpha$  values of the network
- We dropped subscripts on  $a$ , meaning  $a$  is a vector of inputs to the final layer

Because each layer informs the next, we call this a **feedforward** neural network.

## Fitting the model - regression

Training data:  $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$

$x \in \mathbb{R}^p$  ( $p$  features), single output,  $y$

Objective function:

## Fitting the model - regression

Training data:  $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$

$x \in \mathbb{R}^p$  ( $p$  features), single output,  $y$

Objective function:

$$J(\alpha, x, y) = \sum_{i=1}^n (h_{\alpha}(x_i) - y_i)^2 + \lambda \sum_{l=1}^{n_l} \sum_{i=1}^{s_{l+1}} \sum_{j=1}^{s_l} \alpha_{ij}^{(l)}$$

*Handwritten annotations:*  
-  $n_l$ : # layers  
-  $s_{l+1}$ : nodes in  $l+1$  layer  
-  $s_l$ : nodes in  $l$

## Quick notes on objective function and finding parameters

- Form is amenable to classification, just one-hot encode the output and use classification error rate as your objective
- For regression, be sure to scale the *output* variables to lie in the range of the activation function.
  - ▶ For sigmoid, scale to:

## Quick notes on objective function and finding parameters

- Form is amenable to classification, just one-hot encode the output and use classification error rate as your objective
- For regression, be sure to scale the *output* variables to lie in the range of the activation function.
  - ▶ For sigmoid, scale to:  $[0, 1]$
  - ▶ Tanh:

## Quick notes on objective function and finding parameters

- Form is amenable to classification, just one-hot encode the output and use classification error rate as your objective
- For regression, be sure to scale the *output* variables to lie in the range of the activation function.
  - ▶ For sigmoid, scale to:  $[0, 1]$
  - ▶ Tanh:  $[-1, 1]$

## Quick notes on objective function and finding parameters

- Form is amenable to classification, just one-hot encode the output and use classification error rate as your objective
- For regression, be sure to scale the *output* variables to lie in the range of the activation function.
  - ▶ For sigmoid, scale to:  $[0, 1]$
  - ▶ Tanh:  $[-1, 1]$
  - ▶ (I *believe* ReLU requires shifting output to be non-negative; textbook does not address.)
- Solving the objective function involves a form of gradient search
  - ▶ The partial derivatives are found via a technique called backpropagation

## Tensorflow playground

On **this website** you'll find a cool interactive tool that allows you to play with NN for classification.

- ① What are the hyperparameters of the model? Can you explain what each one does?
- ② Try fitting the “exclusive or” (choose on top left) data set.
- ③ Also try fitting the “Spiral” data set.
- ④ Possible spiral solution:



## Tensorflow playground

On **this website** you'll find a cool interactive tool that allows you to play with NN for classification.

- 1 What are the hyperparameters of the model? Can you explain what each one does?
- 2 Try fitting the “exclusive or” (choose on top left) data set.
- 3 Also try fitting the “Spiral” data set.
- 4 Possible spiral solution:
  - 1 Learning rate 0.03
  - 2 Two hidden layers, six and four neurons each
  - 3 Tanh activation
  - 4 Include all but  $X_1X_2$  features.
  - 5 L1 regularization, regularization rate = 0.001

## Tensorflow playground

On **this website** you'll find a cool interactive tool that allows you to play with NN for classification.

- 1 What are the hyperparameters of the model? Can you explain what each one does?
- 2 Try fitting the “exclusive or” (choose on top left) data set.
- 3 Also try fitting the “Spiral” data set.
- 4 Possible spiral solution:
  - 1 Learning rate 0.03
  - 2 Two hidden layers, six and four neurons each
  - 3 Tanh activation
  - 4 Include all but  $X_1X_2$  features.
  - 5 L1 regularization, regularization rate = 0.001
- 5 You got close by trial and error. What's another way?

## Tensorflow playground

On **this website** you'll find a cool interactive tool that allows you to play with NN for classification.

- ① What are the hyperparameters of the model? Can you explain what each one does?
- ② Try fitting the “exclusive or” (choose on top left) data set.
- ③ Also try fitting the “Spiral” data set.
- ④ Possible spiral solution:
  - ① Learning rate 0.03
  - ② Two hidden layers, six and four neurons each
  - ③ Tanh activation
  - ④ Include all but  $X_1X_2$  features.
  - ⑤ L1 regularization, regularization rate = 0.001
- ⑤ You got close by trial and error. What's another way?
  - ▶ Cross validation! Grid search, randomized search
  - ▶ But everything is computationally intense.

## What's going on in the hidden layers?

Hover over the hidden layers in the tensorflow playground.

Q: What are we looking at?

## What's going on in the hidden layers?

Hover over the hidden layers in the tensorflow playground.

Q: What are we looking at?

Ans: The scalar output of that neuron's activation function at each point in the feature space.

These can have interesting (but sometimes dubious) interpretations. More next time!